



VIT[®]
Vellore Institute of Technology
(Deemed to be University under section 3 of UGC Act, 1956)

*****Intrusion Detection System*****

using

*****Machine Learning Algorithms*****

A PROJECT DONE

for

Information Security Management – [F1 slot]

in

B. Tech – Information Technology and Engineering

By

18BIT0126- ROHITH REDDY (narurohith.reddy2018@vitstudent.ac.in)

18BIT0094- SAI SUCHETAN REDDY

SLOT- F1 (BOTH 18BIT0126,0094)

Under the Guidance of

SUMAIYA THASSEN I

isumaiyathassen@vit.ac.in

Associate Professor Grade 1, SITE

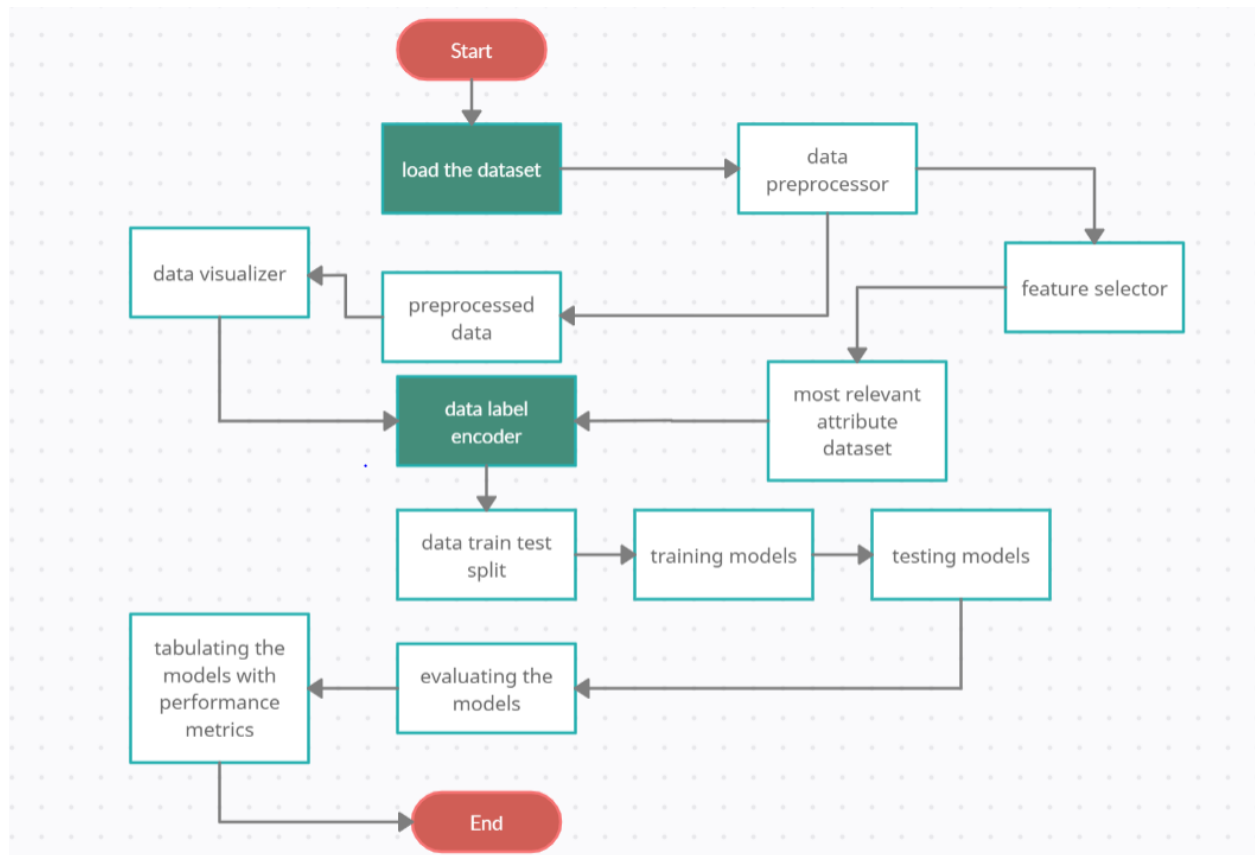
ABSTRACT

Intrusion detection system (IDS) is one of the implemented solutions against harmful attacks. Furthermore, attackers always keep changing their tools and techniques. However, implementing an accepted IDS system is also a challenging task.

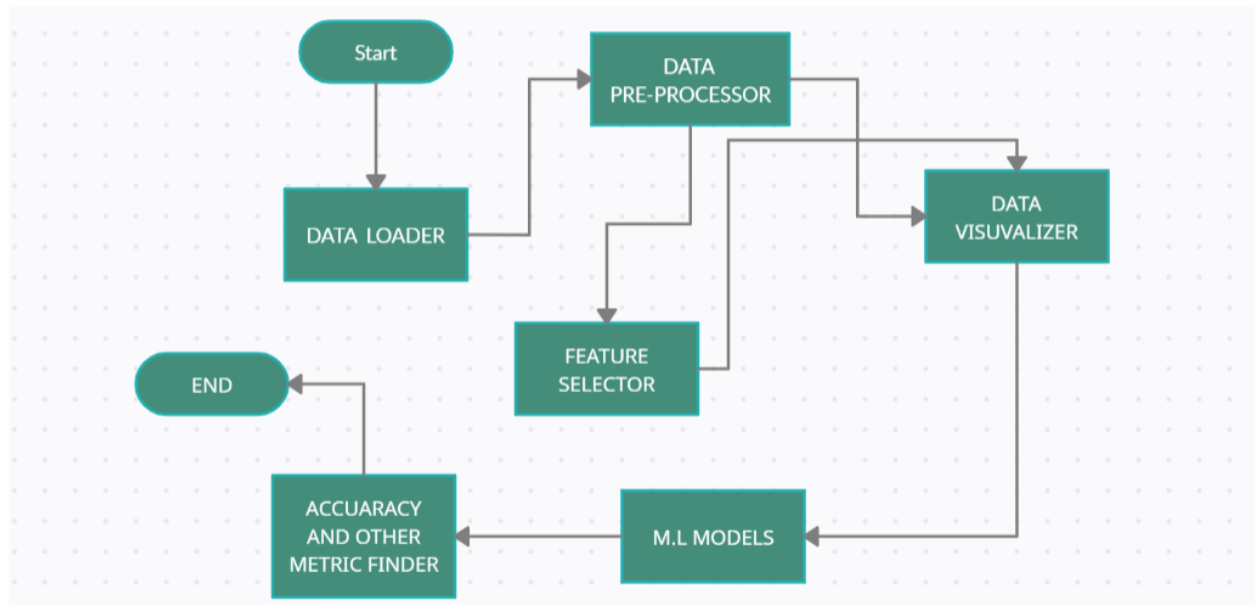
Intrusion Detect Systems (IDSs) are a range of cybersecurity-based technology initially developed to detect vulnerabilities and exploits against a target host. The sole use of the IDS is to identify malicious activities and policy violations. Existing intrusion detection systems rely heavily on human analysts to differentiate intrusive from non-intrusive network traffic. The large and growing amount of data confronts the analysts with an overwhelming task, making the automation of aspects of this task necessary.

In this project, we demonstrate an approach for network Intrusion Detection System (IDS) for cyber security using Machine Learning (ML) techniques using some Supervised Classification Algorithms. And we produce a detail report about the accuracy and the performance of each model we use in our project.

WORKFLOW DIAGRAM



ARCHITECTURE DIAGRAM



intrusion Detection using Machine Learning techniques

18BIT0126

18BIT0094

```
In [1]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn import svm
from sklearn.metrics import classification_report
from sklearn import metrics
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import (precision_score, recall_score, f1_score, accuracy_score, mean_squared_error, mean_absolute_error)
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
```

loading the data from the user machine

```
In [2]: data=pd.read_csv(r'C:\Users\asus\Downloads\NF-UQ-NIDS.csv')
data.describe()
```

Out[2]:

	L4_SRC_PORT	L4_DST_PORT	PROTOCOL	L7_PROTO	IN_BYTES	OUT_BYTES
count	1.199489e+07	1.199489e+07	1.199489e+07	1.199489e+07	1.199489e+07	1.199489e+07
mean	4.124519e+04	8.964758e+03	8.661742e+00	1.932050e+01	3.983727e+03	9.489449e+03
std	2.108654e+04	1.772082e+04	6.404263e+00	3.504857e+01	1.662256e+05	2.933544e+05
min	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00
25%	3.299400e+04	5.300000e+01	6.000000e+00	0.000000e+00	6.800000e+01	0.000000e+00
50%	5.061000e+04	4.430000e+02	6.000000e+00	0.000000e+00	2.320000e+02	1.560000e+02
75%	5.581700e+04	3.389000e+03	6.000000e+00	7.000000e+00	1.440000e+03	1.873000e+03
max	6.553500e+04	6.553500e+04	2.550000e+02	2.510000e+02	2.282235e+08	2.432197e+08

In [3]: data

Out[3]:

	IPV4_SRC_ADDR	L4_SRC_PORT	IPV4_DST_ADDR	L4_DST_PORT	PROTOCOL	L7_PF
0	149.171.126.0	62073	59.166.0.5	56082	6	
1	149.171.126.2	32284	59.166.0.5	1526	6	
2	149.171.126.0	21	59.166.0.1	21971	6	
3	59.166.0.1	23800	149.171.126.0	46893	6	
4	59.166.0.5	63062	149.171.126.2	21	6	
...
11994888	192.168.100.46	80	192.168.100.5	80	6	
11994889	192.168.100.5	0	192.168.100.3	0	6	
11994890	192.168.100.7	365	192.168.100.3	565	17	
11994891	192.168.100.3	50850	13.54.166.67	8883	6	22
11994892	192.168.100.6	49160	192.168.100.149	4444	6	

11994893 rows × 15 columns

data information about datatypes

```
In [4]: import warnings
warnings.filterwarnings('ignore')
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11994893 entries, 0 to 11994892
Data columns (total 15 columns):
IPV4_SRC_ADDR      object
L4_SRC_PORT        int64
IPV4_DST_ADDR      object
L4_DST_PORT        int64
PROTOCOL           int64
L7_PROTO           float64
IN_BYTES           int64
OUT_BYTES          int64
IN_PKTS            int64
OUT_PKTS           int64
TCP_FLAGS          int64
FLOW_DURATION_MILLISECONDS  int64
Label              int64
Attack             object
Dataset            object
dtypes: float64(1), int64(10), object(4)
memory usage: 1.3+ GB
```


dropping unwanted columns

```
In [5]: data=data.drop(['Dataset'],axis=1)
data
```

Out[5]:

	IPV4_SRC_ADDR	L4_SRC_PORT	IPV4_DST_ADDR	L4_DST_PORT	PROTOCOL	L7_PF
0	149.171.126.0	62073	59.166.0.5	56082	6	
1	149.171.126.2	32284	59.166.0.5	1526	6	
2	149.171.126.0	21	59.166.0.1	21971	6	
3	59.166.0.1	23800	149.171.126.0	46893	6	
4	59.166.0.5	63062	149.171.126.2	21	6	
...
11994888	192.168.100.46	80	192.168.100.5	80	6	
11994889	192.168.100.5	0	192.168.100.3	0	6	
11994890	192.168.100.7	365	192.168.100.3	565	17	
11994891	192.168.100.3	50850	13.54.166.67	8883	6	22
11994892	192.168.100.6	49160	192.168.100.149	4444	6	

11994893 rows × 14 columns



```
In [6]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11994893 entries, 0 to 11994892
Data columns (total 14 columns):
IPV4_SRC_ADDR      object
L4_SRC_PORT        int64
IPV4_DST_ADDR      object
L4_DST_PORT        int64
PROTOCOL           int64
L7_PROTO           float64
IN_BYTES           int64
OUT_BYTES          int64
IN_PKTS            int64
OUT_PKTS           int64
TCP_FLAGS          int64
FLOW_DURATION_MILLISECONDS  int64
Label              int64
Attack             object
dtypes: float64(1), int64(10), object(3)
memory usage: 1.3+ GB
```

checking for any null values in the dataset


```
In [7]: data.isnull().sum()
```

```
Out[7]: IPV4_SRC_ADDR      0
L4_SRC_PORT        0
IPV4_DST_ADDR      0
L4_DST_PORT        0
PROTOCOL           0
L7_PROTO           0
IN_BYTES           0
OUT_BYTES          0
IN_PKTS            0
OUT_PKTS           0
TCP_FLAGS          0
FLOW_DURATION_MILLISECONDS  0
Label              0
Attack             0
dtype: int64
```

types of attacks that are in the dataset

```
In [8]: import seaborn as sns
from matplotlib import pyplot as plt
fig, ax=plt.subplots(figsize=(30,30))
sns.countplot(ax=ax,x="Attack", data=data)
```


Out[8]: <matplotlib.axes._subplots.AxesSubplot at 0x21a4ca91bc8>



plotting the labels 0-normal 1-intrusion


```
In [9]: import seaborn as sns
from matplotlib import pyplot as plt
fig, ax=plt.subplots(figsize=(30,30))
sns.countplot(ax=ax,x="Label", data=data)
```

Out[9]: <matplotlib.axes._subplots.AxesSubplot at 0x21a4d25e388>



categorical data into a seperate dataframe

```
In [10]: cdata=data[['IPV4_SRC_ADDR', 'IPV4_DST_ADDR', 'Attack']]
cdata
```

Out[10]:

	IPV4_SRC_ADDR	IPV4_DST_ADDR	Attack
0	149.171.126.0	59.166.0.5	Benign
1	149.171.126.2	59.166.0.5	Benign
2	149.171.126.0	59.166.0.1	Benign
3	59.166.0.1	149.171.126.0	Benign
4	59.166.0.5	149.171.126.2	Benign
...
11994888	192.168.100.46	192.168.100.5	Benign
11994889	192.168.100.5	192.168.100.3	Benign
11994890	192.168.100.7	192.168.100.3	Benign
11994891	192.168.100.3	13.54.166.67	Benign
11994892	192.168.100.6	192.168.100.149	Theft

11994893 rows × 3 columns

label encoding the categorical data

```
In [11]: from sklearn.preprocessing import LabelEncoder
cdata.columns
for label in cdata.columns:
    cdata[label]=LabelEncoder().fit(cdata[label]).transform(cdata[label])
```

In [12]: cdata

Out[12]:

	IPV4_SRC_ADDR	IPV4_DST_ADDR	Attack
0	18035	25333	2
1	18047	25333	2
2	18035	25329	2
3	64615	5246	2
4	64619	5258	2
...
11994888	42254	9060	2
11994889	42255	9057	2
11994890	42258	9057	2
11994891	42252	4508	2
11994892	42257	9054	13

11994893 rows × 3 columns

In [13]: cdata.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11994893 entries, 0 to 11994892
Data columns (total 3 columns):
IPV4_SRC_ADDR    int32
IPV4_DST_ADDR    int32
Attack           int32
dtypes: int32(3)
memory usage: 137.3 MB
```

```
In [14]: data=data.drop(['IPV4_SRC_ADDR','IPV4_DST_ADDR','Attack'],axis=1)
data
```

Out[14]:

	L4_SRC_PORT	L4_DST_PORT	PROTOCOL	L7_PROTO	IN_BYTES	OUT_BYTES	IN_P
0	62073	56082	6	0.000	9672	416	
1	32284	1526	6	0.000	1776	104	
2	21	21971	6	1.000	1842	1236	
3	23800	46893	6	0.000	528	8824	
4	63062	21	6	1.000	1786	2340	
...
11994888	80	80	6	7.000	2330065	0	:
11994889	0	0	6	0.000	1054423	0	
11994890	365	565	17	0.000	62422	0	
11994891	50850	8883	6	222.178	11300	1664	
11994892	49160	4444	6	0.000	40102320	37280	

11994893 rows × 11 columns



```
In [15]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11994893 entries, 0 to 11994892
Data columns (total 11 columns):
L4_SRC_PORT      int64
L4_DST_PORT      int64
PROTOCOL         int64
L7_PROTO         float64
IN_BYTES         int64
OUT_BYTES        int64
IN_PKTS          int64
OUT_PKTS         int64
TCP_FLAGS        int64
FLOW_DURATION_MILLISECONDS int64
Label            int64
dtypes: float64(1), int64(10)
memory usage: 1006.7 MB
```

```
In [16]: fdata = pd.concat([data, cdata], axis=1)
         fdata.head()
```

```
Out[16]:
```

	L4_SRC_PORT	L4_DST_PORT	PROTOCOL	L7_PROTO	IN_BYTES	OUT_BYTES	IN_PKTS	O
0	62073	56082	6	0.0	9672	416	11	
1	32284	1526	6	0.0	1776	104	6	
2	21	21971	6	1.0	1842	1236	26	
3	23800	46893	6	0.0	528	8824	10	
4	63062	21	6	1.0	1786	2340	32	

```
In [17]: fdata.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11994893 entries, 0 to 11994892
Data columns (total 14 columns):
L4_SRC_PORT                int64
L4_DST_PORT                int64
PROTOCOL                  int64
L7_PROTO                  float64
IN_BYTES                  int64
OUT_BYTES                 int64
IN_PKTS                   int64
OUT_PKTS                  int64
TCP_FLAGS                 int64
FLOW_DURATION_MILLISECONDS int64
Label                     int64
IPV4_SRC_ADDR             int32
IPV4_DST_ADDR             int32
Attack                    int32
dtypes: float64(1), int32(3), int64(10)
memory usage: 1.1 GB
```

```
In [18]: x=fdata.drop(['Label'],axis=1)
         y=fdata['Label']
```

train test split

```
In [19]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.30,random_state
              =0)
         x_train.shape,y_train.shape,x_test.shape,y_test.shape
```

```
Out[19]: ((8396425, 13), (8396425,), (3598468, 13), (3598468,))
```

multiple linear regression

```
In [20]: #multiple linear regression
print("***** multiple linear regression*****")
print("*****")
from sklearn.linear_model import LinearRegression
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import accuracy_score, r2_score, mean_squared_error
print("*****making the model....*****")
regressor = LinearRegression()
print("*****training the model....*****")
regressor=regressor.fit(x_train, y_train)
print("*****making prediction please wait....*****")
y_pred = regressor.predict(x_test)
y_pred
print("*****printing confusion matrix and classification report....****")
print("*****")
print(confusion_matrix(y_test, y_pred.round()))
print(classification_report(y_test, y_pred.round()))
```

```
***** multiple linear regression*****
*
*****making the model....*****
*****training the model....*****
*****making prediction please wait....*****
*****printing confusion matrix and classification report....*****
[[      0      0      0      0      0]
 [      5 2762471      14      4      2]
 [     357 178246 621893 35476      0]
 [      0      0      0      0      0]
 [      0      0      0      0      0]]
              precision      recall  f1-score      support

      -1.0      0.00      0.00      0.00      0
      0.0      0.94      1.00      0.97  2762496
      1.0      1.00      0.74      0.85  835972
      2.0      0.00      0.00      0.00      0
      3.0      0.00      0.00      0.00      0

accuracy              0.94  3598468
macro avg      0.39      0.35      0.36  3598468
weighted avg      0.95      0.94      0.94  3598468
```

naive bayes algorithm

```
In [21]: from sklearn import metrics
print("*****the naive bayes algorithm used*****")
gnb=GaussianNB()
print("*****training the model....*****")
gnb.fit(x_train,y_train)
print("*****making the predictions....*****")
y_pred=gnb.predict(x_test)
y_test.value_counts()
print("*****printing confusion matrix and classification report....*****")
cm=confusion_matrix(y_test,y_pred)

print(classification_report(y_test,y_pred))
print(confusion_matrix(y_test,y_pred))
```

```
*****the naive bayes algorithm used*****
*****training the model....*****
*****making the predictions....*****
*****printing confusion matrix and classification report....*****
              precision    recall  f1-score   support

      0       0.94       0.96       0.95     2762496
      1       0.86       0.78       0.82     835972

 accuracy          0.92     3598468
 macro avg         0.90     0.87     0.88     3598468
weighted avg         0.92     0.92     0.92     3598468

[[2658257 104239]
 [ 184495 651477]]
```

decision tree algorithm

```
In [22]: from sklearn import tree
import matplotlib.pyplot as plt
print("*****the decision tree was used*****")
clf=DecisionTreeClassifier(criterion='gini',splitter='random',max_leaf_nodes=1
0,min_samples_leaf=5,max_depth=5,random_state=0)
print("*****training the model....*****")
clf.fit(x_train,y_train)
print("*****making the predictions....*****")
y_pred=clf.predict(x_test)
y_pred
print("*****printing the classification report please wait....*****
*")
print(classification_report(y_test,y_pred))
```

```
*****the decision tree was used*****
*****training the model....*****
*****making the predictions....*****
*****printing the classification report please wait....*****
              precision    recall  f1-score   support

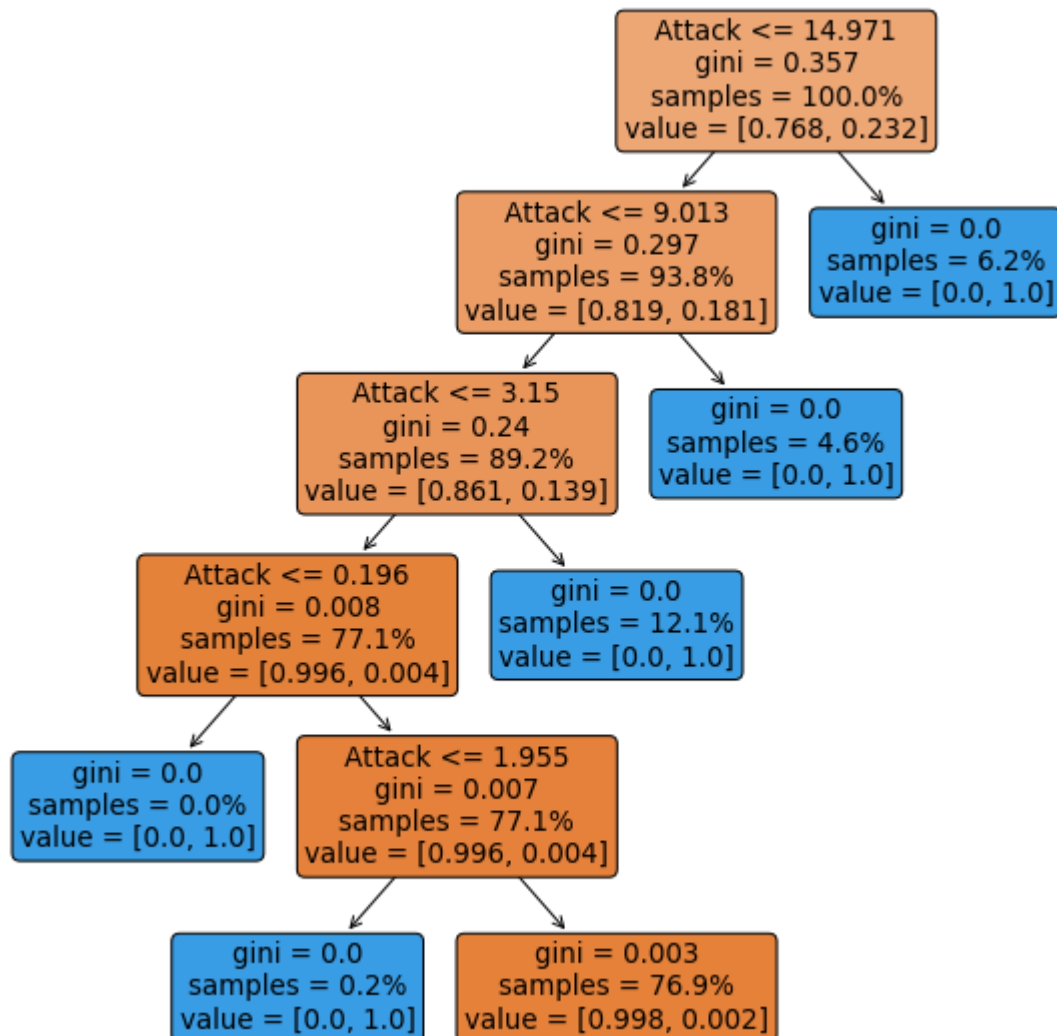
         0           1.00      1.00      1.00     2762496
         1           1.00      0.99      1.00      835972

 accuracy                   1.00     3598468
 macro avg                  1.00      1.00      1.00     3598468
 weighted avg              1.00      1.00      1.00     3598468
```

decision tree visuvalization


```
In [23]: print("*****making the decision tree....*****")
cols=list(x.columns.values)
plt.figure(figsize=(10,10))
tree.plot_tree(clf.fit(x,y),feature_names=cols,filled=True,precision=3,proportion=True,rounded=True)
plt.show()
```

*****making the decision tree....*****



adaboost algorithm

```
In [24]: print("*****the ensemble ADABOOST was used*****")
from sklearn.ensemble import AdaBoostClassifier
model=DecisionTreeClassifier(criterion='entropy',max_depth=1,random_state=0)
print("*****making the model....*****")
Adaboost=AdaBoostClassifier(base_estimator=model,n_estimators=20,random_state=0)
print("*****training the model....*****")
boostmodel=Adaboost.fit(x_train,y_train)
print("*****making the predictions please wait....*****")
y_pred=boostmodel.predict(x_test)
print("*****printing the accuracy of the model please wait....*****")
predictions=metrics.accuracy_score(y_test,y_pred)
print("accuracy: ",predictions)
```

```
*****the ensemble ADABOOST was used*****
*****making the model....*****
*****training the model....*****
*****making the predictions please wait....*****
*****printing the accuracy of the model please wait....*****
accuracy: 1.0
```

random forest algorithm

```
In [25]: #random forest
print("*****the random forest was used*****")
from sklearn.ensemble import RandomForestClassifier
print("*****making the model....*****")
classifier = RandomForestClassifier(n_estimators = 80, criterion = 'entropy',
random_state = 12)
print("*****training the model....*****")
classifier.fit(x_train, y_train)
# predict
print("*****making the predictions....*****")
y_pred = classifier.predict(x_test)
print("*****printing the accuracy....*****")
accuracy = accuracy_score(y_test, y_pred)
print("accuracy: ",accuracy)
```

```
*****the random forest was used*****
*****making the model....*****
*****training the model....*****
*****making the predictions....*****
*****printing the accuracy....*****
accuracy: 1.0
```

logistic regression algorithm

```
In [26]: print("*****the log regression was used*****")
from sklearn.linear_model import LogisticRegression
print("*****making the model....*****")
logmodel = LogisticRegression(max_iter=18000)
print("*****training the model....*****")
logmodel.fit(x_train,y_train)
print("*****making the predictions....*****")
predictions = logmodel.predict(x_test)
print("*****printing classication report and confusion matrix....*****")
print(classification_report(y_test, predictions))
print(confusion_matrix(y_test, predictions))
```

```
*****the log regression was used*****
*****making the model....*****
*****training the model....*****
*****making the predictions....*****
*****printing classication report and confusion matrix....*****
```

	precision	recall	f1-score	support
0	0.93	0.94	0.94	2762496
1	0.80	0.78	0.79	835972
accuracy			0.90	3598468
macro avg	0.87	0.86	0.86	3598468
weighted avg	0.90	0.90	0.90	3598468

```
[[2602019 160477]
 [ 183381 652591]]
```


```
In [ ]: classifier=Sequential()
classifier.add(Dense(output_dim=49,init='he_uniform',activation='relu',input_dim=41))
classifier.add(Dense(output_dim=24,init='he_uniform',activation='relu'))
classifier.add(Dense(output_dim=13,init='he_uniform',activation='relu'))
classifier.add(Dense(output_dim=1,init='glorot_uniforms',activation='sigmoid'))
classifier.compile(optimizer='Adamax',loss='binary_crossentropy',metrics=['accuracy'])
model_history=classifier.fit(x_train,y_train,validation_split=0.33,batch_size=1000,nb_epoch=10)
y_pred=classifier.predict(x_test)
print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test,y_pred))
```

correlation heatmap for feature selection

```

In [27]: import seaborn as sns
X = fdata.drop(['Label'],axis=1) #independent columns
y = fdata['Label'] #target column i.e price range
#get correlations of each features in dataset
corrmat = data.corr()
top_corr_features = corrmat.index
plt.figure(figsize=(20,20))
#plot heat map
g=sns.heatmap(data[top_corr_features].corr(),annot=True,cmap="RdYlGn")

```



```
In [28]: cor=fdata.corr()
cor_target=abs(cor["Label"])
relevant_features=cor_target[cor_target>-1]
relevant_features
```

```
Out[28]: L4_SRC_PORT      0.178658
L4_DST_PORT      0.219108
PROTOCOL        0.140932
L7_PROTO        0.090210
IN_BYTES        0.010478
OUT_BYTES       0.013647
IN_PKTS         0.018832
OUT_PKTS        0.009963
TCP_FLAGS       0.059306
FLOW_DURATION_MILLISECONDS 0.337871
Label          1.000000
IPV4_SRC_ADDR   0.032354
IPV4_DST_ADDR   0.044257
Attack          0.787343
Name: Label, dtype: float64
```

selecting features with correlation>0.3

```
In [29]: ndata=fdata[['FLOW_DURATION_MILLISECONDS', 'Attack']]
ndata
```

```
Out[29]:
```

	FLOW_DURATION_MILLISECONDS	Attack
0	15	2
1	0	2
2	1111	2
3	124	2
4	1459	2
...
11994888	4263037	2
11994889	4263062	2
11994890	4263062	2
11994891	4264935	2
11994892	4270068	13

11994893 rows × 2 columns

```
In [30]: x=ndata
y=fdata['Label']
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.30,random_state=0)
x_train.shape,y_train.shape,x_test.shape,y_test.shape
```

```
Out[30]: ((8396425, 2), (8396425,)), (3598468, 2), (3598468,))
```

```
In [31]: #multiple linear regression
print("***** improved multiple linear regression*****")
print("*****")
from sklearn.linear_model import LinearRegression
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import accuracy_score, r2_score, mean_squared_error
print("*****making the model....*****")
regressor = LinearRegression()
print("*****training the model....*****")
regressor=regressor.fit(x_train, y_train)
print("*****making prediction please wait....*****")
y_pred = regressor.predict(x_test)
y_pred
print("*****printing confusion matrix and classification report....*****")
print("*****")
print(confusion_matrix(y_test, y_pred.round()))
print(classification_report(y_test, y_pred.round()))
```

```
***** improved multiple linear regression*****
*****
*****making the model....*****
*****training the model....*****
*****making prediction please wait....*****
*****printing confusion matrix and classification report....*****
[[2762496      0      0]
 [ 432728 366815 36429]
 [      0      0      0]]
              precision    recall  f1-score   support

         0.0         0.86      1.00      0.93    2762496
         1.0         1.00      0.44      0.61    835972
         2.0         0.00      0.00      0.00         0

 accuracy          0.87    3598468
 macro avg         0.62    0.48    0.51    3598468
 weighted avg         0.90    0.87    0.85    3598468
```

```
In [32]: from sklearn import metrics
print("*****the improved naive bayes algorithm used*****")
)
gnb=GaussianNB()
print("*****training the model....*****")
gnb.fit(x_train,y_train)
print("*****making the predictions....*****")
y_pred=gnb.predict(x_test)
y_test.value_counts()
print("*****printing confusion matrix and classification report....*****")
cm=confusion_matrix(y_test,y_pred)

print(classification_report(y_test,y_pred))
print(confusion_matrix(y_test,y_pred))
```

```
*****the improved naive bayes algorithm used*****
*****training the model....*****
*****making the predictions....*****
*****printing confusion matrix and classification report....*****
```

	precision	recall	f1-score	support
0	0.77	1.00	0.87	2762496
1	0.00	0.00	0.00	835972
accuracy			0.77	3598468
macro avg	0.38	0.50	0.43	3598468
weighted avg	0.59	0.77	0.67	3598468

```
[[2762496    0]
 [ 835972    0]]
```

```
In [33]: print("*****the modified log regression was used*****")
from sklearn.linear_model import LogisticRegression
print("*****making the model....*****")
logmodel = LogisticRegression(max_iter=18000)
print("*****training the model....*****")
logmodel.fit(x_train,y_train)
print("*****making the predictions....*****")
predictions = logmodel.predict(x_test)
print("*****printing classication report and confusion matrix....*****")
print(classification_report(y_test, predictions))
print(confusion_matrix(y_test, predictions))
```

```
*****the modified log regression was used*****
*****making the model....*****
*****training the model....*****
*****making the predictions....*****
*****printing classication report and confusion matrix....*****
```

	precision	recall	f1-score	support
0	0.82	0.73	0.77	2762496
1	0.35	0.47	0.40	835972
accuracy			0.67	3598468
macro avg	0.58	0.60	0.59	3598468
weighted avg	0.71	0.67	0.69	3598468

```
[[2014931  747565]
 [ 439370  396602]]
```

In []:

The above was the over all model that we have made rather than a single machine learning model running for IDS. We felt that two models out of which one predicts the attack and the other does the prediction of the intrusion which help us to provide extra layer of security and we can even stop the intrusion before the attack is made by those predictions.

The processing time of algorithms that we had used

[11994893 rows almost 1.3 Gb data for predictions]

Multiple linear regression acc=94	32 sec
Naïve bayes algorithm acc=92	30 sec
Decision tree algorithm acc=1	34 sec
Ada-boost algorithm acc=1	8-10 min
Random forest algorithm acc=1	10-14 min
Logistic regression algorithm acc=90	5-8 min
Hyper parameter tuned logistic regression acc=1	1-2 hrs

We had calculated the ROC-AUC curves for the prediction models that are attached below

After that the attack prediction model and then the label prediction model are attached

ROC and AUC curves for the models developed above

```
In [1]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn import svm
from sklearn.metrics import classification_report
from sklearn import metrics
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import (precision_score, recall_score, f1_score, accuracy_score, mean_squared_error, mean_absolute_error)
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
```

```
In [2]: data=pd.read_csv(r'C:\Users\asus\Downloads\NF-UQ-NIDS.csv')
data.describe()
```

Out[2]:

	L4_SRC_PORT	L4_DST_PORT	PROTOCOL	L7_PROTO	IN_BYTES	OUT_BYTES
count	1.199489e+07	1.199489e+07	1.199489e+07	1.199489e+07	1.199489e+07	1.199489e+07
mean	4.124519e+04	8.964758e+03	8.661742e+00	1.932050e+01	3.983727e+03	9.489449e+03
std	2.108654e+04	1.772082e+04	6.404263e+00	3.504857e+01	1.662256e+05	2.933544e+05
min	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00
25%	3.299400e+04	5.300000e+01	6.000000e+00	0.000000e+00	6.800000e+01	0.000000e+00
50%	5.061000e+04	4.430000e+02	6.000000e+00	0.000000e+00	2.320000e+02	1.560000e+02
75%	5.581700e+04	3.389000e+03	6.000000e+00	7.000000e+00	1.440000e+03	1.873000e+03
max	6.553500e+04	6.553500e+04	2.550000e+02	2.510000e+02	2.282235e+08	2.432197e+08

```
In [3]: data
```

Out[3]:

	IPV4_SRC_ADDR	L4_SRC_PORT	IPV4_DST_ADDR	L4_DST_PORT	PROTOCOL	L7_PF
0	149.171.126.0	62073	59.166.0.5	56082	6	
1	149.171.126.2	32284	59.166.0.5	1526	6	
2	149.171.126.0	21	59.166.0.1	21971	6	
3	59.166.0.1	23800	149.171.126.0	46893	6	
4	59.166.0.5	63062	149.171.126.2	21	6	
...
11994888	192.168.100.46	80	192.168.100.5	80	6	
11994889	192.168.100.5	0	192.168.100.3	0	6	
11994890	192.168.100.7	365	192.168.100.3	565	17	
11994891	192.168.100.3	50850	13.54.166.67	8883	6	22
11994892	192.168.100.6	49160	192.168.100.149	4444	6	

11994893 rows × 15 columns

data information about datatypes

```
In [4]: import warnings
warnings.filterwarnings('ignore')
data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11994893 entries, 0 to 11994892
Data columns (total 15 columns):
IPV4_SRC_ADDR          object
L4_SRC_PORT            int64
IPV4_DST_ADDR          object
L4_DST_PORT            int64
PROTOCOL               int64
L7_PROTO               float64
IN_BYTES               int64
OUT_BYTES              int64
IN_PKTS                int64
OUT_PKTS               int64
TCP_FLAGS              int64
FLOW_DURATION_MILLISECONDS int64
Label                  int64
Attack                 object
Dataset                object
dtypes: float64(1), int64(10), object(4)
memory usage: 1.3+ GB
```

dropping unwanted columns

```
In [5]: data=data.drop(['Dataset'],axis=1)
```

categorical data into a seperate dataframe

```
In [6]: cdata=data[['IPV4_SRC_ADDR', 'IPV4_DST_ADDR', 'Attack']]
cdata
```

Out[6]:

	IPV4_SRC_ADDR	IPV4_DST_ADDR	Attack
0	149.171.126.0	59.166.0.5	Benign
1	149.171.126.2	59.166.0.5	Benign
2	149.171.126.0	59.166.0.1	Benign
3	59.166.0.1	149.171.126.0	Benign
4	59.166.0.5	149.171.126.2	Benign
...
11994888	192.168.100.46	192.168.100.5	Benign
11994889	192.168.100.5	192.168.100.3	Benign
11994890	192.168.100.7	192.168.100.3	Benign
11994891	192.168.100.3	13.54.166.67	Benign
11994892	192.168.100.6	192.168.100.149	Theft

11994893 rows × 3 columns

label encoding the categorical data


```
In [7]: from sklearn.preprocessing import LabelEncoder
cdata.columns
for label in cdata.columns:
    cdata[label]=LabelEncoder().fit(cdata[label]).transform(cdata[label])
```

```
In [8]: data=data.drop(['IPV4_SRC_ADDR','IPV4_DST_ADDR','Attack'],axis=1)
data
```

Out[8]:

	L4_SRC_PORT	L4_DST_PORT	PROTOCOL	L7_PROTO	IN_BYTES	OUT_BYTES	IN_P
0	62073	56082	6	0.000	9672	416	
1	32284	1526	6	0.000	1776	104	
2	21	21971	6	1.000	1842	1236	
3	23800	46893	6	0.000	528	8824	
4	63062	21	6	1.000	1786	2340	
...
11994888	80	80	6	7.000	2330065	0	:
11994889	0	0	6	0.000	1054423	0	
11994890	365	565	17	0.000	62422	0	
11994891	50850	8883	6	222.178	11300	1664	
11994892	49160	4444	6	0.000	40102320	37280	


11994893 rows × 11 columns



```
In [9]: fdata = pd.concat([data, cdata], axis=1)
fdata.head()
```

Out[9]:

	L4_SRC_PORT	L4_DST_PORT	PROTOCOL	L7_PROTO	IN_BYTES	OUT_BYTES	IN_PKTS	O
0	62073	56082	6	0.0	9672	416	11	
1	32284	1526	6	0.0	1776	104	6	
2	21	21971	6	1.0	1842	1236	26	
3	23800	46893	6	0.0	528	8824	10	
4	63062	21	6	1.0	1786	2340	32	



```
In [10]: fdata.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11994893 entries, 0 to 11994892
Data columns (total 14 columns):
L4_SRC_PORT                int64
L4_DST_PORT                int64
PROTOCOL                   int64
L7_PROTO                   float64
IN_BYTES                   int64
OUT_BYTES                  int64
IN_PKTS                    int64
OUT_PKTS                   int64
TCP_FLAGS                  int64
FLOW_DURATION_MILLISECONDS int64
Label                      int64
IPV4_SRC_ADDR              int32
IPV4_DST_ADDR              int32
Attack                     int32
dtypes: float64(1), int32(3), int64(10)
memory usage: 1.1 GB
```

```
In [11]: x=fdata.drop(['Label'],axis=1)
         y=fdata['Label']
```

train test split

```
In [12]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.30,random_state
           =0)
         x_train.shape,y_train.shape,x_test.shape,y_test.shape
```

```
Out[12]: ((8396425, 13), (8396425,), (3598468, 13), (3598468,))
```

naive bayes algorithm

```
In [13]: from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import accuracy_score, r2_score, mean_squared_error
from sklearn import metrics
print("*****the naive bayes algorithm used*****")
gnb=GaussianNB()
print("*****training the model....*****")
gnb.fit(x_train,y_train)
print("*****making the predictions....*****")
y_pred=gnb.predict(x_test)
y_test.value_counts()
print("*****printing confusion matrix and classification report....*****")
cm=confusion_matrix(y_test,y_pred)

print(classification_report(y_test,y_pred))
print(confusion_matrix(y_test,y_pred))
```

```
*****the naive bayes algorithm used*****
*****training the model....*****
*****making the predictions....*****
*****printing confusion matrix and classification report....*****
              precision    recall  f1-score   support

     0       0.94      0.96      0.95     2762496
     1       0.86      0.78      0.82     835972

 accuracy          0.92     3598468
 macro avg       0.90      0.87      0.88     3598468
 weighted avg    0.92      0.92      0.92     3598468


[[2658257 104239]
 [ 184495 651477]]
```



```
In [29]: # roc curve and auc
from sklearn.datasets import make_classification
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from matplotlib import pyplot

# generate a no skill prediction (majority class)
ns_probs = [0 for _ in range(len(y_test))]
print("predicting the probabilities")
# predict probabilities
lr_probs = gnb.predict_proba(x_test)
# keep probabilities for the positive outcome only
lr_probs = lr_probs[:, 1]
print("calculating scores")
# calculate scores
ns_auc = roc_auc_score(y_test, ns_probs)
lr_auc = roc_auc_score(y_test, lr_probs)
print("summerizing the scores")
# summarize scores
print('No Skill: ROC AUC=%.3f' % (ns_auc))
print('naive bayesian: ROC AUC=%.3f' % (lr_auc))
print("calculating roc curve")
# calculate roc curves
ns_fpr, ns_tpr, _ = roc_curve(y_test, ns_probs)
lr_fpr, lr_tpr, _ = roc_curve(y_test, lr_probs)
print("plotting roc curve")
# plot the roc curve for the model
pyplot.plot(ns_fpr, ns_tpr, linestyle='--', label='No Skill')
pyplot.plot(lr_fpr, lr_tpr, marker='.', label='Logistic')
# axis labels
pyplot.xlabel('False Positive Rate')
pyplot.ylabel('True Positive Rate')
# show the legend
pyplot.legend()
# show the plot
pyplot.show()
```

predicting the probabilities
 calculating scores
 summerizing the scores
 No Skill: ROC AUC=0.500
 naive bayesan: ROC AUC=0.927
 calculating roc curve
 plotting roc curve



decision tree algorithm

```

In [17]: from sklearn import tree
import matplotlib.pyplot as plt
print("*****the decision tree was used*****")
clf=DecisionTreeClassifier(criterion='gini',splitter='random',max_leaf_nodes=1
0,min_samples_leaf=5,max_depth=5,random_state=0)
print("*****training the model....*****")
clf.fit(x_train,y_train)
print("*****making the predictions....*****")
y_pred=clf.predict(x_test)
y_pred
print("*****printing the classification report please wait....*****
*")
print(classification_report(y_test,y_pred))
  
```

```

*****the decision tree was used*****
*****training the model....*****
*****making the predictions....*****
*****printing the classification report please wait....*****
              precision    recall  f1-score   support


     0           1.00       1.00       1.00     2762496
     1           1.00       0.99       1.00       835972

 accuracy          1.00          1.00          1.00     3598468
 macro avg          1.00          1.00          1.00     3598468
 weighted avg       1.00          1.00          1.00     3598468
  
```

```
In [28]: # generate a no skill prediction (majority class)
ns_probs = [0 for _ in range(len(y_test))]

print("predicting the probabilities")
# predict probabilities
lr_probs = clf.predict_proba(x_test)
# keep probabilities for the positive outcome only
lr_probs = lr_probs[:, 1]
print("calculating scores")
# calculate scores
ns_auc = roc_auc_score(y_test, ns_probs)
lr_auc = roc_auc_score(y_test, lr_probs)
print("summerizing the scores")
# summarize scores
print('No Skill: ROC AUC=%.3f' % (ns_auc))
print('decision tree: ROC AUC=%.3f' % (lr_auc))
print("calculating roc curve")
# calculate roc curves
ns_fpr, ns_tpr, _ = roc_curve(y_test, ns_probs)
lr_fpr, lr_tpr, _ = roc_curve(y_test, lr_probs)
print("plotting roc curve")
# plot the roc curve for the model
pyplot.plot(ns_fpr, ns_tpr, linestyle='--', label='No Skill')
pyplot.plot(lr_fpr, lr_tpr, marker='.', label='Logistic')
# axis labels
pyplot.xlabel('False Positive Rate')
pyplot.ylabel('True Positive Rate')
# show the legend
pyplot.legend()
# show the plot
pyplot.show()
```

predicting the probabilities
calculating scores
summerizing the scores
No Skill: ROC AUC=0.500
decision tree: ROC AUC=0.997
calculating roc curve
plotting roc curve



adaboost algorithm

```
In [19]: print("*****the ensemble ADABOOST was used*****")
from sklearn.ensemble import AdaBoostClassifier
model=DecisionTreeClassifier(criterion='entropy',max_depth=1,random_state=0)
print("*****making the model....*****")
Adaboost=AdaBoostClassifier(base_estimator=model,n_estimators=20,random_state=
0)
print("*****training the model....*****")
boostmodel=Adaboost.fit(x_train,y_train)
print("*****making the predictions please wait....*****")
y_pred=boostmodel.predict(x_test)
print("*****printing the accuracy of the model please wait....*****
*")
predictions=metrics.accuracy_score(y_test,y_pred)
print("accuracy: ",predictions)
```

```
*****the ensemble ADABOOST was used*****
*****making the model....*****
*****training the model....*****
*****making the predictions please wait....*****
*****printing the accuracy of the model please wait....*****
accuracy:  1.0
```


```

In [27]: # generate a no skill prediction (majority class)
ns_probs = [0 for _ in range(len(y_test))]
# fit a model

print("predicting the probabilities")
# predict probabilities
lr_probs = boostmodel.predict_proba(x_test)
# keep probabilities for the positive outcome only
lr_probs = lr_probs[:, 1]
print("calculating scores")
# calculate scores
ns_auc = roc_auc_score(y_test, ns_probs)
lr_auc = roc_auc_score(y_test, lr_probs)
print("summerizing the scores")
# summarize scores
print('No Skill: ROC AUC=%.3f' % (ns_auc))
print('adaboost: ROC AUC=%.3f' % (lr_auc))
print("calculating roc curve")
# calculate roc curves
ns_fpr, ns_tpr, _ = roc_curve(y_test, ns_probs)
lr_fpr, lr_tpr, _ = roc_curve(y_test, lr_probs)
print("plotting roc curve")
# plot the roc curve for the model
pyplot.plot(ns_fpr, ns_tpr, linestyle='--', label='No Skill')
pyplot.plot(lr_fpr, lr_tpr, marker='.', label='Logistic')
# axis labels
pyplot.xlabel('False Positive Rate')
pyplot.ylabel('True Positive Rate')
# show the legend
pyplot.legend()
# show the plot
pyplot.show()

```

predicting the probabilities
 calculating scores
 summerizing the scores
 No Skill: ROC AUC=0.500
 adaboost: ROC AUC=1.000
 calculating roc curve
 plotting roc curve



```
In [21]: #random forest
print("*****the random forest was used*****")
from sklearn.ensemble import RandomForestClassifier
print("*****making the model....*****")
classifier = RandomForestClassifier(n_estimators = 80, criterion = 'entropy',
random_state = 12)
print("*****training the model....*****")
classifier.fit(x_train, y_train)
# predict
print("*****making the predictions....*****")
y_pred = classifier.predict(x_test)
print("*****printing the accuracy....*****")
accuracy = accuracy_score(y_test, y_pred)
print("accuracy: ",accuracy)
```


```
*****the random forest was used*****
*****making the model....*****
*****training the model....*****
*****making the predictions....*****
*****printing the accuracy....*****
accuracy:  1.0
```

```

In [26]: # generate a no skill prediction (majority class)
ns_probs = [0 for _ in range(len(y_test))]
print("predicting the probabilities")
# predict probabilities
lr_probs = classifier.predict_proba(x_test)
# keep probabilities for the positive outcome only
lr_probs = lr_probs[:, 1]
print("calculating scores")
# calculate scores
ns_auc = roc_auc_score(y_test, ns_probs)
lr_auc = roc_auc_score(y_test, lr_probs)
print("summerizing the scores")
# summarize scores
print('No Skill: ROC AUC=%.3f' % (ns_auc))
print('random forest: ROC AUC=%.3f' % (lr_auc))
print("calculating roc curve")
# calculate roc curves
ns_fpr, ns_tpr, _ = roc_curve(y_test, ns_probs)
lr_fpr, lr_tpr, _ = roc_curve(y_test, lr_probs)
print("plotting roc curve")
# plot the roc curve for the model
pyplot.plot(ns_fpr, ns_tpr, linestyle='--', label='No Skill')
pyplot.plot(lr_fpr, lr_tpr, marker='.', label='Logistic')
# axis labels
pyplot.xlabel('False Positive Rate')
pyplot.ylabel('True Positive Rate')
# show the legend
pyplot.legend()
# show the plot
pyplot.show()

```

predicting the probabilities
 calculating scores
 summerizing the scores
 No Skill: ROC AUC=0.500
 random forest: ROC AUC=1.000
 calculating roc curve
 plotting roc curve



logistic regression algorithm

```
In [24]: print("*****the log regression was used*****")
from sklearn.linear_model import LogisticRegression
print("*****making the model....*****")
logmodel = LogisticRegression(max_iter=18000)
print("*****training the model....*****")
logmodel.fit(x_train,y_train)
print("*****making the predictions....*****")
predictions = logmodel.predict(x_test)
print("*****printing classication report and confusion matrix....*****")
print(classification_report(y_test, predictions))
print(confusion_matrix(y_test, predictions))
```

```
*****the log regression was used*****
*****making the model....*****
*****training the model....*****
*****making the predictions....*****
*****printing classication report and confusion matrix....*****
```

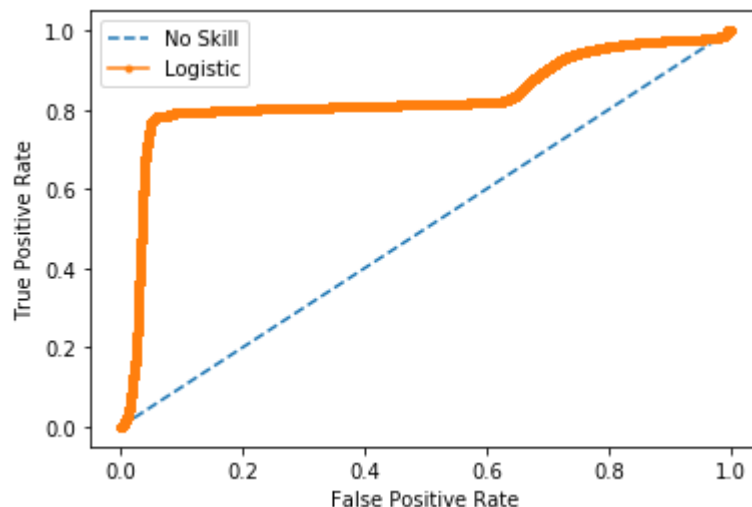
	precision	recall	f1-score	support
0	0.93	0.94	0.94	2762496
1	0.80	0.78	0.79	835972
accuracy			0.90	3598468
macro avg	0.87	0.86	0.86	3598468
weighted avg	0.90	0.90	0.90	3598468

```
[[2602019 160477]
 [ 183381 652591]]
```



```
In [25]: # generate a no skill prediction (majority class)
ns_probs = [0 for _ in range(len(y_test))]
print("predicting the probabilities")
# predict probabilities
lr_probs = logmodel.predict_proba(x_test)
# keep probabilities for the positive outcome only
lr_probs = lr_probs[:, 1]
print("calculating scores")
# calculate scores
ns_auc = roc_auc_score(y_test, ns_probs)
lr_auc = roc_auc_score(y_test, lr_probs)
print("summerizing the scores")
# summarize scores
print('No Skill: ROC AUC=%.3f' % (ns_auc))
print('logistic: ROC AUC=%.3f' % (lr_auc))
print("calculating roc curve")
# calculate roc curves
ns_fpr, ns_tpr, _ = roc_curve(y_test, ns_probs)
lr_fpr, lr_tpr, _ = roc_curve(y_test, lr_probs)
print("plotting roc curve")
# plot the roc curve for the model
pyplot.plot(ns_fpr, ns_tpr, linestyle='--', label='No Skill')
pyplot.plot(lr_fpr, lr_tpr, marker='.', label='Logistic')
# axis labels
pyplot.xlabel('False Positive Rate')
pyplot.ylabel('True Positive Rate')
# show the legend
pyplot.legend()
# show the plot
pyplot.show()
```

predicting the probabilities
calculating scores
summerizing the scores
No Skill: ROC AUC=0.500
logistic: ROC AUC=0.831
calculating roc curve
plotting roc curve



**attack
prediction
model**


```
In [1]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn import svm
from sklearn.metrics import classification_report
from sklearn import metrics
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import (precision_score, recall_score, f1_score, accuracy_score, mean_squared_error, mean_absolute_error)
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
```

```
In [2]: data=pd.read_csv(r'C:\Users\asus\Downloads\NF-UQ-NIDS.csv')
data.describe()
data=data.drop(['Dataset'],axis=1)
data=data.drop(['Label'],axis=1)
data
```

Out[2]:

	IPV4_SRC_ADDR	L4_SRC_PORT	IPV4_DST_ADDR	L4_DST_PORT	PROTOCOL	L7_PF
0	149.171.126.0	62073	59.166.0.5	56082	6	
1	149.171.126.2	32284	59.166.0.5	1526	6	
2	149.171.126.0	21	59.166.0.1	21971	6	
3	59.166.0.1	23800	149.171.126.0	46893	6	
4	59.166.0.5	63062	149.171.126.2	21	6	
...
11994888	192.168.100.46	80	192.168.100.5	80	6	
11994889	192.168.100.5	0	192.168.100.3	0	6	
11994890	192.168.100.7	365	192.168.100.3	565	17	
11994891	192.168.100.3	50850	13.54.166.67	8883	6	22
11994892	192.168.100.6	49160	192.168.100.149	4444	6	

11994893 rows × 13 columns



```
In [3]: import warnings
warnings.filterwarnings('ignore')
numda=data[['Attack']]
numda.head()
```

Out[3]:

	Attack
0	Benign
1	Benign
2	Benign
3	Benign
4	Benign

```
In [4]: cdata=data[['IPV4_SRC_ADDR', 'IPV4_DST_ADDR', 'Attack']]
cdata
```

Out[4]:

	IPV4_SRC_ADDR	IPV4_DST_ADDR	Attack
0	149.171.126.0	59.166.0.5	Benign
1	149.171.126.2	59.166.0.5	Benign
2	149.171.126.0	59.166.0.1	Benign
3	59.166.0.1	149.171.126.0	Benign
4	59.166.0.5	149.171.126.2	Benign
...
11994888	192.168.100.46	192.168.100.5	Benign
11994889	192.168.100.5	192.168.100.3	Benign
11994890	192.168.100.7	192.168.100.3	Benign
11994891	192.168.100.3	13.54.166.67	Benign
11994892	192.168.100.6	192.168.100.149	Theft

11994893 rows × 3 columns


```
In [5]: from sklearn.preprocessing import LabelEncoder
cdata.columns
for label in cdata.columns:
    cdata[label]=LabelEncoder().fit(cdata[label]).transform(cdata[label])
```

```
In [6]: data=data.drop(['IPV4_SRC_ADDR','IPV4_DST_ADDR','Attack'],axis=1)
data
```

Out[6]:

	L4_SRC_PORT	L4_DST_PORT	PROTOCOL	L7_PROTO	IN_BYTES	OUT_BYTES	IN_P
0	62073	56082	6	0.000	9672	416	
1	32284	1526	6	0.000	1776	104	
2	21	21971	6	1.000	1842	1236	
3	23800	46893	6	0.000	528	8824	
4	63062	21	6	1.000	1786	2340	
...
11994888	80	80	6	7.000	2330065	0	
11994889	0	0	6	0.000	1054423	0	
11994890	365	565	17	0.000	62422	0	
11994891	50850	8883	6	222.178	11300	1664	
11994892	49160	4444	6	0.000	40102320	37280	

11994893 rows × 10 columns



```
In [7]: fdata = pd.concat([data, cdata], axis=1)
fdata.head()
```

Out[7]:

	L4_SRC_PORT	L4_DST_PORT	PROTOCOL	L7_PROTO	IN_BYTES	OUT_BYTES	IN_PKTS	O
0	62073	56082	6	0.0	9672	416	11	
1	32284	1526	6	0.0	1776	104	6	
2	21	21971	6	1.0	1842	1236	26	
3	23800	46893	6	0.0	528	8824	10	
4	63062	21	6	1.0	1786	2340	32	



```
In [8]: numda['numdata']=fdata[['Attack']]
numda.head()
```

Out[8]:

	Attack	numdata
0	Benign	2
1	Benign	2
2	Benign	2
3	Benign	2
4	Benign	2

```
In [9]: numda = numda.drop_duplicates('Attack')
print(numda)
```

	Attack	numdata
0	Benign	2
29	Exploits	7
67	Reconnaissance	11
93	DoS	6
548	Generic	9
600	Shellcode	12
753	Backdoor	1
1858	Fuzzers	8
2663	Worms	14
35786	Analysis	0
1623130	injection	15
1623131	DDoS	5
1623136	scanning	19
1623264	password	17
1624507	mitm	16
1816782	xss	20
2403109	ransomware	18
3042976	Infiltration	10
4620416	Bot	3
4919431	Brute Force	4
11394794	Theft	13

```
In [10]: x=fdata.drop(['Attack'],axis=1)
y=fdata['Attack']
```

```
In [11]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.30,random_state
=0)
x_train.shape,y_train.shape,x_test.shape,y_test.shape
```

```
Out[11]: ((8396425, 12), (8396425,), (3598468, 12), (3598468,))
```

```
In [12]: print("*****the ensemble ADABOOST was used*****")
from sklearn.ensemble import AdaBoostClassifier
model=DecisionTreeClassifier(criterion='entropy',max_depth=1,random_state=0)
print("*****making the model....*****")
Adaboost=AdaBoostClassifier(base_estimator=model,n_estimators=20,random_state=
0)
print("*****training the model....*****")
boostmodel=Adaboost.fit(x_train,y_train)
print("*****making the predictions please wait....*****")
y_pred=boostmodel.predict(x_test)
print("*****printing the accuracy of the model please wait....*****
*")
predictions=metrics.accuracy_score(y_test,y_pred)
print("accuracy: ",predictions)

*****the ensemble ADABOOST was used*****
*****making the model....*****
*****training the model....*****
*****making the predictions please wait....*****
*****printing the accuracy of the model please wait....*****
accuracy:  0.6075938427130657
```

```
In [13]: #random forest
print("*****the random forest was used*****")
from sklearn.ensemble import RandomForestClassifier
print("*****making the model....*****")
classifier = RandomForestClassifier(n_estimators = 80, criterion = 'entropy',
random_state = 12)
print("*****training the model....*****")
classifier.fit(x_train, y_train)
# predict
print("*****making the predictions....*****")
y_pred = classifier.predict(x_test)
print("*****printing the accuracy....*****")
accuracy = accuracy_score(y_test, y_pred)
print("accuracy: ",accuracy)

*****the random forest was used*****
*****making the model....*****
*****training the model....*****
*****making the predictions....*****
*****printing the accuracy....*****
accuracy:  0.9320055090110569
```

label

prediction

model



```
In [1]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn import svm
from sklearn.metrics import classification_report
from sklearn import metrics
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import (precision_score, recall_score, f1_score, accuracy_score, mean_squared_error, mean_absolute_error)
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
```

```
In [2]: data=pd.read_csv(r'C:\Users\asus\Downloads\NF-UQ-NIDS.csv')
data.describe()
data=data.drop(['Dataset'],axis=1)
data
```

Out[2]:

	IPV4_SRC_ADDR	L4_SRC_PORT	IPV4_DST_ADDR	L4_DST_PORT	PROTOCOL	L7_PF
0	149.171.126.0	62073	59.166.0.5	56082	6	
1	149.171.126.2	32284	59.166.0.5	1526	6	
2	149.171.126.0	21	59.166.0.1	21971	6	
3	59.166.0.1	23800	149.171.126.0	46893	6	
4	59.166.0.5	63062	149.171.126.2	21	6	
...
11994888	192.168.100.46	80	192.168.100.5	80	6	
11994889	192.168.100.5	0	192.168.100.3	0	6	
11994890	192.168.100.7	365	192.168.100.3	565	17	
11994891	192.168.100.3	50850	13.54.166.67	8883	6	22
11994892	192.168.100.6	49160	192.168.100.149	4444	6	

11994893 rows × 14 columns



```
In [3]: import warnings
warnings.filterwarnings('ignore')
data=data.drop(['Attack'],axis=1)
```

```
In [4]: cdata=data[['IPV4_SRC_ADDR', 'IPV4_DST_ADDR']]
cdata
```

Out[4]:

	IPV4_SRC_ADDR	IPV4_DST_ADDR
0	149.171.126.0	59.166.0.5
1	149.171.126.2	59.166.0.5
2	149.171.126.0	59.166.0.1
3	59.166.0.1	149.171.126.0
4	59.166.0.5	149.171.126.2
...
11994888	192.168.100.46	192.168.100.5
11994889	192.168.100.5	192.168.100.3
11994890	192.168.100.7	192.168.100.3
11994891	192.168.100.3	13.54.166.67
11994892	192.168.100.6	192.168.100.149

11994893 rows × 2 columns

```
In [5]: from sklearn.preprocessing import LabelEncoder
cdata.columns
for label in cdata.columns:
    cdata[label]=LabelEncoder().fit(cdata[label]).transform(cdata[label])
```

```
In [6]: data=data.drop(['IPV4_SRC_ADDR', 'IPV4_DST_ADDR'],axis=1)
data
```

Out[6]:


	L4_SRC_PORT	L4_DST_PORT	PROTOCOL	L7_PROTO	IN_BYTES	OUT_BYTES	IN_P
0	62073	56082	6	0.000	9672	416	
1	32284	1526	6	0.000	1776	104	
2	21	21971	6	1.000	1842	1236	
3	23800	46893	6	0.000	528	8824	
4	63062	21	6	1.000	1786	2340	
...	
11994888	80	80	6	7.000	2330065	0	;
11994889	0	0	6	0.000	1054423	0	
11994890	365	565	17	0.000	62422	0	
11994891	50850	8883	6	222.178	11300	1664	
11994892	49160	4444	6	0.000	40102320	37280	

11994893 rows × 11 columns

```
In [7]: fdata = pd.concat([data, cdata], axis=1)
        fdata.head()
```

Out[7]:

	L4_SRC_PORT	L4_DST_PORT	PROTOCOL	L7_PROTO	IN_BYTES	OUT_BYTES	IN_PKTS	O
0	62073	56082	6	0.0	9672	416	11	
1	32284	1526	6	0.0	1776	104	6	
2	21	21971	6	1.0	1842	1236	26	
3	23800	46893	6	0.0	528	8824	10	
4	63062	21	6	1.0	1786	2340	32	



```
In [8]: x=fdata.drop(['Label'],axis=1)
        y=fdata['Label']
```

```
In [9]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.30,random_state=0)
        x_train.shape,y_train.shape,x_test.shape,y_test.shape
```

Out[9]: ((8396425, 12), (8396425,), (3598468, 12), (3598468,))

```
In [10]: #multiple linear regression
print("***** multiple linear regression*****")
print("*****")
from sklearn.linear_model import LinearRegression
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import accuracy_score, r2_score, mean_squared_error
print("*****making the model....*****")
regressor = LinearRegression()
print("*****training the model....*****")
regressor=regressor.fit(x_train, y_train)
print("*****making prediction please wait....*****")
y_pred = regressor.predict(x_test)
y_pred
print("*****printing confusion matrix and classification report....*****")
print("*****")
print(confusion_matrix(y_test, y_pred.round()))
print(classification_report(y_test, y_pred.round()))
```

```
***** multiple linear regression*****
*
*****making the model....*****
*****training the model....*****
*****making prediction please wait....*****
*****printing confusion matrix and classification report....*****
[[      0      0      0      0      0      0      0      0]
 [      0      0      0      0      0      0      0      0]
 [      0      0      0      0      0      0      0      0]
 [      0      0      0      0      0      0      0      0]
 [      0      0      4  12794 2584447 165246      4      1]
 [      1      1   948   1568  234021  599433      0      0]
 [      0      0      0      0      0      0      0      0]
 [      0      0      0      0      0      0      0      0]]

              precision    recall  f1-score   support

-5.0              0.00        0.00        0.00         0
-4.0              0.00        0.00        0.00         0
-2.0              0.00        0.00        0.00         0
-1.0              0.00        0.00        0.00         0
 0.0              0.92        0.94        0.93    2762496
 1.0              0.78        0.72        0.75    835972
 2.0              0.00        0.00        0.00         0
 3.0              0.00        0.00        0.00         0

 accuracy              0.88    3598468
 macro avg              0.21        0.21        0.21    3598468
 weighted avg           0.89        0.88        0.89    3598468
```

```
In [11]: from sklearn import metrics
print("*****the naive bayes algorithm used*****")
gnb=GaussianNB()
print("*****training the model....*****")
gnb.fit(x_train,y_train)
print("*****making the predictions....*****")
y_pred=gnb.predict(x_test)
y_test.value_counts()
print("*****printing confusion matrix and classification report....*****")
cm=confusion_matrix(y_test,y_pred)

print(classification_report(y_test,y_pred))
print(confusion_matrix(y_test,y_pred))
```

```
*****the naive bayes algorithm used*****
*****training the model....*****
*****making the predictions....*****
*****printing confusion matrix and classification report....*****
```

	precision	recall	f1-score	support
0	0.93	0.96	0.95	2762496
1	0.86	0.78	0.82	835972
accuracy			0.92	3598468
macro avg	0.90	0.87	0.88	3598468
weighted avg	0.92	0.92	0.92	3598468

```
[[2657889 104607]
 [ 184952 651020]]
```

```
In [12]: from sklearn import tree
import matplotlib.pyplot as plt
print("*****the decision tree was used*****")
clf=DecisionTreeClassifier(criterion='gini',splitter='random',max_leaf_nodes=1
0,min_samples_leaf=5,max_depth=5,random_state=0)
print("*****training the model....*****")
clf.fit(x_train,y_train)
print("*****making the predictions....*****")
y_pred=clf.predict(x_test)
y_pred
print("*****printing the classification report please wait....*****
*")
print(classification_report(y_test,y_pred))
```


```
*****the decision tree was used*****
*****training the model....*****
*****making the predictions....*****
*****printing the classification report please wait....*****
              precision    recall  f1-score   support

         0           0.98       0.96       0.97    2762496
         1           0.88       0.92       0.90     835972

 accuracy                   0.95    3598468
 macro avg           0.93       0.94       0.93    3598468
 weighted avg        0.95       0.95       0.95    3598468
```

```
In [13]: print("*****making the decision tree....*****")
cols=list(x.columns.values)
plt.figure(figsize=(10,10))
tree.plot_tree(clf.fit(x,y),feature_names=cols,filled=True,precision=3,proportion=True,rounded=True)
plt.show()
```

*****making the decision tree....*****



```
In [14]: print("*****the ensemble ADABOOST was used*****")
from sklearn.ensemble import AdaBoostClassifier
model=DecisionTreeClassifier(criterion='entropy',max_depth=1,random_state=0)
print("*****making the model....*****")
Adaboost=AdaBoostClassifier(base_estimator=model,n_estimators=20,random_state=
0)
print("*****training the model....*****")
boostmodel=Adaboost.fit(x_train,y_train)
print("*****making the predictions please wait....*****")
y_pred=boostmodel.predict(x_test)
print("*****printing the accuracy of the model please wait....*****
*")
predictions=metrics.accuracy_score(y_test,y_pred)
print("accuracy: ",predictions)

*****the ensemble ADABOOST was used*****
*****making the model....*****
*****training the model....*****
*****making the predictions please wait....*****
*****printing the accuracy of the model please wait....*****
accuracy:  0.9831884012863252
```

```
In [15]: #random forest
print("*****the random forest was used*****")
from sklearn.ensemble import RandomForestClassifier
print("*****making the model....*****")
classifier = RandomForestClassifier(n_estimators = 80, criterion = 'entropy',
random_state = 12)
print("*****training the model....*****")
classifier.fit(x_train, y_train)
# predict
print("*****making the predictions....*****")
y_pred = classifier.predict(x_test)
print("*****printing the accuracy....*****")
accuracy = accuracy_score(y_test, y_pred)
print("accuracy: ",accuracy)

*****the random forest was used*****
*****making the model....*****
*****training the model....*****
*****making the predictions....*****
*****printing the accuracy....*****
accuracy:  0.9948883802773847
```



```
In [16]: print("*****the log regression was used*****")
from sklearn.linear_model import LogisticRegression
print("*****making the model....*****")
logmodel = LogisticRegression(max_iter=18000)
print("*****training the model....*****")
logmodel.fit(x_train,y_train)
print("*****making the predictions....*****")
predictions = logmodel.predict(x_test)
print("*****printing classication report and confusion matrix....*****")
print(classification_report(y_test, predictions))
print(confusion_matrix(y_test, predictions))
```

```
*****the log regression was used*****
*****making the model....*****
*****training the model....*****
*****making the predictions....*****
*****printing classication report and confusion matrix....*****
```

	precision	recall	f1-score	support
0	0.93	0.94	0.94	2762496
1	0.80	0.78	0.79	835972
accuracy			0.90	3598468
macro avg	0.87	0.86	0.86	3598468
weighted avg	0.90	0.90	0.90	3598468

```
[[2602005 160491]
 [ 183376 652596]]
```

```
In [ ]: from sklearn.model_selection import GridSearchCV
logistic = LogisticRegression(max_iter=30000)
penalty = [ 'l2' ]
C = np.logspace(0, 4, 10)
hyperparameters = dict(C=C, penalty=penalty)
clf = GridSearchCV(logistic, hyperparameters, cv=5, verbose=0)
best_model = clf.fit(x_train, y_train)
print('Best Penalty:', best_model.best_estimator_.get_params()['penalty'])
print('Best C:', best_model.best_estimator_.get_params()['C'])
```

```
*****the log regression was used*****
*****making the model....*****
*****training the model....*****
*****making the predictions....*****
*****printing classication report and confusion matrix....*****
precision    recall  f1-score   support

      0         1.00      1.00      1.00     2762496
      1         1.00      1.00      1.00     835972

accuracy              1.00     3598468
macro avg           1.00      1.00      1.00     3598468
weighted avg       1.00      1.00      1.00     3598468

[[2602005  160491]
 [ 183376  652596]]
```

In []: